

# Fault Tolerant Distributed Consensus

There is no conversation more boring than the one where everybody agrees.

Marco Aiello, Eirini Kaldeli

University of Groningen

Distributed Systems, 2009

# Chapter 14: Consensus and Agreement

Ajay Kshemkalyani and Mukesh Singhal


Distributed Computing: Principles, Algorithms, and Systems

Cambridge University Press

# Assumptions

## System assumptions

- Failure models
- Synchronous/ Asynchronous communication
- Network connectivity
- Sender identification
- Channel reliability
- Authenticated vs. non-authenticated messages
- Agreement variable



# Problem Specifications

Byzantine Agreement (single source has an initial value)

**Agreement:** All non-faulty processes must agree on the same value.

**Validity:** If the source process is non-faulty, then the agreed upon value by all the non-faulty processes must be the same as the initial value of the source.

**Termination:** Each non-faulty process must eventually decide on a value.

Consensus Problem (all processes have an initial value)

**Agreement:** All non-faulty processes must agree on the same (single) value.

**Validity:** If all the non-faulty processes have the same initial value, then the agreed upon value by all the non-faulty processes must be that same value.

**Termination:** Each non-faulty process must eventually decide on a value.

Interactive Consistency (all processes have an initial value)

**Agreement:** All non-faulty processes must agree on the same array of values  $A[v_1 \dots v_n]$ .

**Validity:** If process  $i$  is non-faulty and its initial value is  $v_i$ , then all non-faulty processes agree on  $v_i$  as the  $i$ th element of the array  $A$ . If process  $j$  is faulty, then the non-faulty processes can agree on any value for  $A[j]$ .

**Termination:** Each non-faulty process must eventually decide on the array  $A$ .

These problems are equivalent to one another! Show using reductions.

# Overview of Results

Failure mode	Synchronous system (message-passing and shared memory)	Asynchronous system (message-passing and shared memory)
No failure	agreement attainable; common knowledge also attainable	agreement attainable; concurrent common knowledge attainable
Crash failure	agreement attainable $f < n$ Byzantine processes $\Omega(f + 1)$ rounds	agreement not attainable
Byzantine failure	agreement attainable $f \leq \lfloor (n - 1)/3 \rfloor$ Byzantine processes $\Omega(f + 1)$ rounds	agreement not attainable

**Table:** Overview of results on agreement.  $f$  denotes number of failure-prone processes.  $n$  is the total number of processes.

In a failure-free system, consensus can be attained in a straightforward manner

## Formal requirements for the consensus algorithm

An algorithm solves the consensus problem if it satisfies the following formal properties:

- **Termination:** Eventually every **non-faulty** processor decides on a value  $y_i$ .
- **Agreement:** The final decisions of all **non-faulty** processors are identical, i.e. if  $y_i, y_j$  are assigned then  $\forall p_i, p_j \in \text{Nonfaulty} : (y_i = y_j)$ .
- **Validity:** If all non-faulty  $p_i$ s have the same input then the decision of a **non-faulty** processor equals the common input, i.e. if  $\forall p_i \in \text{Nonfaulty} : (x_i = v)$  then if  $y_i$  is assigned for some non-faulty  $p_j$  then  $y_j = v$ .

Systems with different levels of synchrony or different kinds of failures require different algorithms.

# Impossibility of Distributed Consensus with one faulty process

## [Fisher, Lynch and Peterson, 1985]

### Bad news

The design of a consensus protocol that tolerates failures is impossible in asynchronous distributed systems.

- Impossibility holds for both shared memory and message passing systems:
  - even if we assume reliable communication channels.
  - even if considering only benign failures (crashes).
  - even if at most one processor fails.
- The problem is that in totally asynchronous systems we cannot distinguish a dead process from a merely slow one.
- We have to assume some level of synchrony on communication, processes or message order for consensus to become possible.

# Impossibility Result (MP, async)

## FLP Impossibility result

Impossible to reach consensus in an async MP system even if a single process has a crash failure

- In a failure-free async MP system, initial state is *monovalent*  $\implies$  consensus can be reached.
- In the face of failures, initial state is necessarily bivalent
- Transforming the input assignments from the all-0 case to the all-1 case, there must exist input assignments  $\vec{I}_a$  and  $\vec{I}_b$  that are 0-valent and 1-valent, resp., and that differ in the input value of only one process, say  $P_i$ . If a 1-failure tolerant consensus protocol exists, then:
  - ▶ Starting from  $\vec{I}_a$ , if  $P_i$  fails immediately, the other processes must agree on 0 due to the termination condition.
  - ▶ Starting from  $\vec{I}_b$ , if  $P_i$  fails immediately, the other processes must agree on 1 due to the termination condition.

However, execution (2) looks identical to execution (1), to all processes, and must end with a consensus value of 0, a contradiction. Hence, there must exist at least one bivalent initial state.

- Consensus requires some communication of initial values.



# Impossibility Result (MP, async)

- To transition from bivalent to monovalent step, must exist a critical step which allows the transition by making a decision
- Critical step cannot be local (cannot tell apart between slow and failed process) nor can it be across multiple processes (it would not be well-defined)
- Hence, cannot transit from bivalent to univalent state.

## Wider Significance of Impossibility Result

- By showing reduction from consensus to problem X, then X is also not solvable under same model (single crash failure)
- E.g., leader election, terminating reliable broadcast, atomic broadcast, computing a network-wide global function using BC-CC flows, transaction commit.

Fault-tolerant consensus can be reached in synchronous systems under certain assumptions on the number of faulty processors and the connectivity of the communication graph.

We will consider a model that can accommodate for process failures:

- The system includes at most  $f$  faulty processors: *f-resilient*.
- The subset  $F$  of the faulty processors (maybe different in each execution) is not known in advance.
- Communication channels are reliable (compare with the two generals problem).
- The graph topology is a complete graph.

## A solely crash tolerant consensus algorithm

Code for each processor  $p_i$ ,  $1 \leq i \leq n$ :


```
V = {  $x_i$  }           /*set V contains  $p_i$ 's input*/  
for (k := 1 to f+1)   /*round k*/  
  broadcast ( $u \in V$  :  $p_i$  has not already sent  $u$ )  
  receive set of msgs  $S_j$  from  $p_j$ ,  $1 \leq j \leq n$ ,  $j \neq i$   
   $V := V \cup \bigcup_{j=1}^n S_j$  /*update V by joining it with the received sets*/  
  
 $y_i = \text{majority}(V)$  /*decide at  $f + 1$  round*/
```

Correctness of the algorithm:

- *Termination*: The algorithm requires exactly  $f + 1$  rounds.
- *Validity*: The decision value is an input of some  $p_i$ , since no spurious messages are introduced: if all inputs have the same value, then that is the only one ever in circulation.

*Agreement:* At the end of round  $f + 1$ ,  
 $\forall p_i, p_j \notin F : (x \in V_i \Rightarrow x \in V_j)$ : prove by contradiction.

## Proof.

- Suppose  $\exists x : (x \in V_i) \wedge (x \notin V_j)$ , where  $p_i, p_j$  non-faulty.
- $p_i$  must have received  $x$  for the **first** time at round  $f + 1$ , otherwise it would have already sent it to  $p_j$ .
- There is a  $p_{i_{f+1}}$  that sent  $x$  to  $p_i$  at round  $f + 1$ .  $p_{i_{f+1}}$  must have crashed in middle of this round, so  $x$  was not sent to  $p_j$ .
- Similarly, there is a  $p_{i_f}$  that sent  $x$  to  $p_{i_{f+1}}$ .
- So, there is a chain of  $f + 1$  distinct faulty processors  $p_{i_1}, \dots, p_{i_{f+1}}$  (remember that after a crash there is no resurrection), that transferred  $x$  to  $p_i$   **Contradiction**.



- Number of messages sent:  $O(n^2)$ . There are at most  $n$  different values and each of them is sent at most  $n - 1$  times.
- Number of rounds:  $f + 1$
- It can be proved that  $f + 1$  is the **lower bound on rounds** for reaching fault-tolerant consensus (both for the benign and the severe case).
- Note that the algorithm is correct no matter how many the faulty processors are.

- There are  $n$  generals who head different divisions of the Byzantine army and have to agree whether to attack the enemy or not.
- Communication is reliable but  $f$  of the generals are traitors and try to bring confusion by feeding incorrect information.
- How many traitors can a byzantine consensus protocol tolerate?


### Theorem

*In a system with three processors one of which are byzantine, there is no algorithm that solves the consensus problem.*

☺ Let's see why.

# Upper Bound on Byzantine Processes (sync)

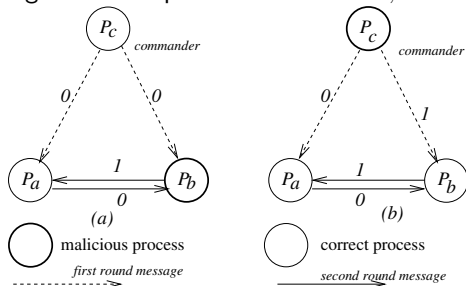
Agreement impossible when  $f = 1, n = 3$ .



- Taking simple majority decision does not help because loyal commander  $P_c$  cannot distinguish between the possible scenarios (a) and (b);
- hence does not know which action to take.
- Proof using induction that problem solvable if  $f \leq \lfloor \frac{n-1}{3} \rfloor$ . See text.

# Upper Bound on Byzantine Processes (sync)

Agreement impossible when  $f = 1, n = 3$ .



- Taking simple majority decision does not help because loyal commander  $P_c$  cannot distinguish between the possible scenarios (a) and (b);
- hence does not know which action to take.
- Proof using induction that problem solvable if  $f \leq \lfloor \frac{n-1}{3} \rfloor$ . See text.



# Lower bound on the ratio of faulty processors to achieve byzantine consensus

## Theorem


*In a system with  $n$  processors  $f$  of which are byzantine, there is no algorithm that solves the consensus problem if  $n \leq 3f$  (even if the network is synchronous and complete).*

We can show that if we assume that there is such an algorithm, then we would be able to solve the problem for  $n = 3$ ,  $f = 1$  contradicting the previous theorem.

- What if the graph is not complete?
- It can be proved that for byzantine consensus to be possible the connectivity of the graph has to be at least  $2f + 1$ .

- ✓ We will present a byzantine consensus algorithm that is optimal in terms of resilience ( $f < \frac{n}{3}$ ) and number of rounds ( $f+1$ ).
- ✗ However, the size of the exchanged messages is exponential.
- Each  $p_i$  maintains a tree data structure of height  $f + 1$  (levels 0 to  $f + 1$ ).
- The algorithm consists of two phases:
  - ① Information gathering: Values are filled in the tree level by level during the  $f + 1$  rounds.
  - ② Decision phase: Each  $p_i$  calculates its decision based on the values in its tree.

# Consensus Solvable when $f = 1, n = 4$



- There is no ambiguity at any loyal commander, when taking majority decision
- Majority decision is over 2nd round messages, and 1st round message received directly from commander-in-chief process.

# Byzantine Generals (recursive formulation), (sync, msg-passing)

(variables)

**boolean:**  $v \leftarrow$  initial value;

**integer:**  $f \leftarrow$  maximum number of malicious processes,  $\leq \lfloor (n - 1)/3 \rfloor$ ;

(message type)

$Oral\_Msg(v, Dests, List, faulty)$ , where

$v$  is a boolean,

$Dests$  is a set of destination process ids to which the message is sent,

$List$  is a list of process ids traversed by this message, ordered from most recent to earliest,

$faulty$  is an integer indicating the number of malicious processes to be tolerated.

$Oral\_Msg(f)$ , where  $f > 0$ :

- 1 The algorithm is initiated by the Commander, who sends his source value  $v$  to all other processes using a  $OM(v, N, \langle i \rangle, f)$  message. The commander returns his own value  $v$  and terminates.
- 2 **[Recursion unfolding:]** For each message of the form  $OM(v_j, Dests, List, f')$  received in this round from some process  $j$ , the process  $i$  uses the value  $v_j$  it receives from the source, and using that value, acts as a new source. (If no value is received, a default value is assumed.)  
To act as a new source, the process  $i$  initiates  $Oral\_Msg(f' - 1)$ , wherein it sends  $OM(v_j, Dests - \{i\}, concat(\langle i \rangle, L), (f' - 1))$  to destinations not in  $concat(\langle i \rangle, L)$  in the next round.
- 3 **[Recursion folding:]** For each message of the form  $OM(v_j, Dests, List, f')$  received in Step 2, each process  $i$  has computed the agreement value  $v_k$ , for each  $k$  not in  $List$  and  $k \neq i$ , corresponding to the value received from  $P_k$  after traversing the nodes in  $List$ , at one level lower in the recursion. If it receives no value in this round, it uses a default value. Process  $i$  then uses the value  $majority_{k \notin List, k \neq i}(v_j, v_k)$  as the agreement value and returns it to the next higher level in the recursive invocation.

$Oral\_Msg(0)$ :

- 1 **[Recursion unfolding:]** Process acts as a source and sends its value to each other process.
- 2 **[Recursion folding:]** Each process uses the value it receives from the other sources, and uses that value as the agreement value. If no value is received, a default value is assumed.

# Relationship between # Messages and Rounds

round number	a message has already visited	aims to tolerate these many failures	and each message gets sent to	total number of messages in round
1	1	$f$	$n - 1$	$n - 1$
2	2	$f - 1$	$n - 2$	$(n - 1) \cdot (n - 2)$
...	...	...	...	...
$x$	$x$	$(f + 1) - x$	$n - x$	$(n - 1)(n - 2) \dots (n - x)$
$x + 1$	$x + 1$	$(f + 1) - x - 1$	$n - x - 1$	$(n - 1)(n - 2) \dots (n - x - 1)$
$f + 1$	$f + 1$	0	$n - f - 1$	$(n - 1)(n - 2) \dots (n - f - 1)$

**Table:** Relationships between messages and rounds in the Oral Messages algorithm for Byzantine agreement.

Complexity:  $f + 1$  rounds, exponential amount of space, and

$$(n - 1) + (n - 1)(n - 2) + \dots + (n - 1)(n - 2) \dots (n - f - 1) \text{ messages}$$

# Bzantine Generals (iterative formulation), Sync, Msg-passing

(variables)

**boolean:**  $v \leftarrow$  initial value;

**integer:**  $f \leftarrow$  maximum number of malicious processes,  $\leq \lfloor \frac{n-1}{3} \rfloor$ ;

**tree of boolean:**

- level 0 root is  $v_{init}^L$ , where  $L = \langle \rangle$ ;

- level  $h (f \geq h > 0)$  nodes: for each  $v_j^L$  at level  $h - 1 = \text{sizeof}(L)$ , its  $n - 2 - \text{sizeof}(L)$  descendants at level  $h$  are  $v_k^{\text{concat}(\langle j \rangle, L)}$ ,  $\forall k$  such that  $k \neq j$ ,  $i$  and  $k$  is not a member of list  $L$ .

(message type)

$OM(v, Dests, List, faulty)$ , where the parameters are as in the recursive formulation.

(1) Initiator (i.e., Commander) initiates Oral Byzantine agreement:

(1a) **send**  $OM(v, N - \{i\}, \langle P_i \rangle, f)$  to  $N - \{i\}$ ;

(1b) **return**( $v$ ).

(2) (Non-initiator, i.e., Lieutenant) receives Oral Message  $OM$ :

(2a) **for**  $rnd = 0$  **to**  $f$  **do**

(2b) **for** each message  $OM$  that arrives in this round, **do**

(2c) **receive**  $OM(v, Dests, L = \langle P_{k_1} \dots P_{k_{f+1-faulty}} \rangle, faulty)$  from  $P_{k_1}$ ;  
 $// \text{ faulty} + \text{round} = f; |Dests| + \text{sizeof}(L) = n$

(2d)  $v_{\text{head}(L)}^{\text{tail}(L)} \leftarrow v$ ;  $// \text{ sizeof}(L) + \text{faulty} = f + 1$ . fill in estimate.


(2e) **send**  $OM(v, Dests - \{i\}, \langle P_i, P_{k_1} \dots P_{k_{f+1-faulty}} \rangle, faulty - 1)$  to  $Dests - \{i\}$  if  $rnd < f$ ;

(2f) **for**  $level = f - 1$  **down to**  $0$  **do**

(2g) **for** each of the  $1 \cdot (n - 2) \cdot \dots \cdot (n - (level + 1))$  nodes  $v_x^L$  in level  $level$ , **do**

(2h)  $v_x^L(x \neq i, x \notin L) = \text{majority}_{y \notin \text{concat}(\langle x \rangle, L); y \neq i}(v_x^L, v_y^{\text{concat}(\langle x \rangle, L)})$ ;

# Tree Data Structure for Agreement Problem (Byzantine Generals)



Some branches of the tree at  $P_3$ . In

this example,  $n = 10$ ,  $f = 3$ , commander is  $P_0$ .

- (round 1)  $P_0$  sends its value to all other processes using  $Oral\_Msg(3)$ , including to  $P_3$ .
- (round 2)  $P_3$  sends 8 messages to others (excl.  $P_0$  and  $P_3$ ) using  $Oral\_Msg(2)$ .  $P_3$  also receives 8 messages.
- (round 3)  $P_3$  sends  $8 \times 7 = 56$  messages to all others using  $Oral\_Msg(1)$ ;  $P_3$  also receives 56 messages.
- (round 4)  $P_3$  sends  $56 \times 6 = 336$  messages to all others using  $Oral\_Msg(0)$ ;  $P_3$  also receives 336 messages. The received values are used as estimates of the majority function at this level of recursion.

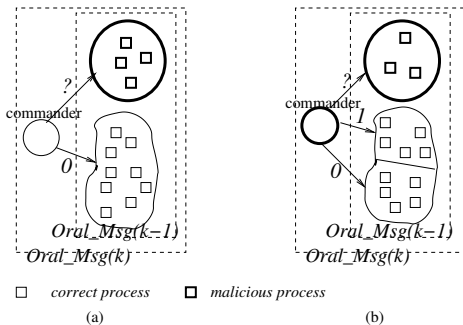
# Exponential Algorithm: An example

An example of the majority computation is as follows.

- $P_3$  revises its estimate of  $v_7^{(5,0)}$  by taking  $\text{majority}(v_7^{(5,0)}, v_1^{(7,5,0)}, v_2^{(7,5,0)}, v_4^{(7,5,0)}, v_6^{(7,5,0)}, v_8^{(7,5,0)}, v_9^{(7,5,0)})$ . Similarly for the other nodes at level 2 of the tree.
- $P_3$  revises its estimate of  $v_5^{(0)}$  by taking  $\text{majority}(v_5^{(0)}, v_1^{(5,0)}, v_2^{(5,0)}, v_4^{(5,0)}, v_6^{(5,0)}, v_7^{(5,0)}, v_8^{(5,0)}, v_9^{(5,0)})$ . Similarly for the other nodes at level 1 of the tree.
- $P_3$  revises its estimate of  $v_0^{(\cdot)}$  by taking  $\text{majority}(v_0^{(\cdot)}, v_1^{(0)}, v_2^{(0)}, v_4^{(0)}, v_5^{(0)}, v_6^{(0)}, v_7^{(0)}, v_8^{(0)}, v_9^{(0)})$ . This is the consensus value.



# Impact of a Loyal and of a Disloyal Commander




The effects of a loyal or a disloyal commander in a system with  $n = 14$  and  $f = 4$ . The subsystems that need to tolerate  $k$  and  $k - 1$  traitors are shown for two cases. (a) Loyal commander. (b) No assumptions about commander.

(a) the commander who invokes  $Oral\_Msg(x)$  is loyal, so all the loyal processes have the same estimate. Although the subsystem of  $3x$  processes has  $x$  malicious processes, all the loyal processes have the same view to begin with. Even if this case repeats for each nested invocation of  $Oral\_Msg$ , even after  $x$  rounds, among the processes, the loyal processes are in a simple majority, so the majority function works in having them maintain the same common view of the loyal commander's value.

(b) the commander who invokes  $Oral\_Msg(x)$  may be malicious and can send conflicting values to the loyal processes. The subsystem of  $3x$  processes has  $x - 1$  malicious processes, but all the loyal processes do not have the same view to begin with.

## EIG algorithm: The tree structure

The EIG tree data structure with  $n = 4$  and  $f = 1$ : the labels of nodes




## The EIG algorithm: Information gathering

- Initially each  $p_i$  stores its input in the root (level 0).
- Round  $1 \leq r \leq f + 1$ , each  $p_i$ :
  - broadcasts all nodes of the  $r - 1$ th level of its tree.
  - fills in level  $r$ : when it receives a message from  $p_j$  with the value of the node labeled  $v = i_1, \dots, i_k$ , it stores it to the node labeled  $v, j$  of its tree (if a value for a node is not received, then default  $u_{\perp}$  is stored).
- So,  $p_i$  stores in node  $i_1, \dots, i_k, j$  the value that “ $p_j$  says that  $p_{i_k}$  says that ... that  $p_{i_1}$  said”.

## The EIG algorithm: Information gathering

The tree is filled in from the root to the leaves, level by level.  
Information received from  $p_3$  in round 2 is stored at level 2:




## The EIG algorithm: Decision phase

- At round  $f + 1$  the entire tree has been filled in. Node labeled with sequence  $\pi$  has value  $tree(\pi)$ .
- Each  $p_i$  applies to each subtree with root  $\pi$  a recursive reduction function (usually majority vote)  $resolve_i(\pi)$ .
- The decision value is the resolved value of the root,  $resolve()$ , which is computed recursively based on the following definition:

$$resolve(\pi) = \begin{cases} tree(\pi) & \text{if } \pi \text{ is a leaf} \\ majority\{resolve(\pi'), \pi' : \text{child of } \pi\} & \text{otherwise} \\ (u_{\perp} & \text{if no majority exists}) \end{cases}$$


## The EIG algorithm: Decision phase

Start from the leaves, and compute the resolved value level by level till the root. Assume binary consensus (input  $\in \{0, 1\}$ ). Default value is 0.




## The EIG algorithm: Decision phase

Start from the leaves, and compute the resolved value level by level till the root. Assume binary consensus (input  $\in \{0, 1\}$ ). Default value is 0.



## The EIG algorithm: Decision phase

Start from the leaves, and compute the resolved value level by level till the root. Assume binary consensus (input  $\in \{0, 1\}$ ). Default value is 0.





Validity can be proved based on the following lemma that says that the resolved values are consistent.

### Lemma 1

For each non-faulty  $p_i$ , the resolved value of a node  $\pi = \pi'j$  corresponding to a non-faulty  $p_j$  is the value that  $p_j$  had stored in its  $\pi'$  node:  $resolve_i(\pi) = tree_j(\pi')$ .

- ☺ Let's prove it (induction on the height of the tree).
- Suppose all non-faulty  $p_i$ s start with initial value  $v$ .
  - The decision value of  $p_i$  is  
 $resolve_i() = majority\{resolve_i(j), j : \text{child of root}\}$
  - By lemma 1 we have that  
 $\forall \text{ non-faulty } j : (resolve_i(j) = tree_j() = v)$ .
  - Since the majority of  $p_i$ s are non-faulty  $resolve_i() = v$ .

To prove agreement we need to introduce some additional terms.

- A node  $\pi$  is **common** if  $\forall p_i, p_j \in \text{Nonfaulty} : (\text{resolve}_i(\pi) = \text{resolve}_j(\pi))$ .
- a node  $\pi$  has a **common frontier** if every path from  $\pi$  to a leaf contains a common node.

### Lemma 2

If a node  $\pi$  has a common frontier then  $\pi$  is common.

☺ Let's prove it (induction on height of the tree, by contradiction).

Agreement can be proven based on lemma 2.

- The nodes on each path from a node at level 1 to a leaf correspond to  $f + 1$  different processors.
- So, at least one such node  $\pi$  corresponds to a non-faulty  $p_j$ , and by lemma 1 its resolved values are consistent (equal to  $tree_j(\pi')$ , where  $\pi = \pi'j$ ), and thus it is common.
- Thus, the root has a common frontier since every path from the root to the leaves includes a common node.
- By lemma 2 the root is common, meaning that all non-faulty processors resolve the same decision value.

The EIG algorithm uses:


- $f + 1$  rounds (optimal)
- $n \geq 3f + 1$  processors (optimal)
- exponential size messages (sub-optimal):
  - At each round  $r$  every process broadcasts the whole level  $r$  of its tree.
  - At  $r = 1$  each  $p_i$  broadcasts one value, at  $r = 2$   $n$  values, at  $r = 3$   $n(n - 1)$  values and at  $r = k$   $n(n - 1)(n - 2) \dots (n - (r - 2))$  values.
  - The largest message corresponds to  $r = f + 1$ :  
 $n(n - 1) \dots (n - (f + 1)) = O(n^f)$ .
- $n^2(f + 1)$  number of messages (sub-optimal)

- ✓ We will present an algorithm that uses message of  $O(1)$  size.
- ✗ However, at the cost of  $2(f + 1)$  rounds and with the requirement that  $n > 4f$ .
- The algorithm contains  $f + 1$  phases, each taking 2 rounds.
- Each  $p_i$  has a *preference* at each phase, which is initially its input value and becomes its decision at phase  $f + 1$ .
- At each phase  $k$   $p_k$  is said to be the *king* of the phase.

# The Phase King Algorithm

## Operation

- Each round has a unique "phases king" derived, say, from PID.
- Each round has two phases:
  - in 1st phase, each process sends its estimate to all other processes.
  - in 2nd phase, the "Phase king" process arrives at an estimate based on the values it received in 1st phase, and broadcasts its new estimate to all others.



# The Phase King Algorithm: Code

(variables)

**boolean:**  $v \leftarrow$  initial value;

**integer:**  $f \leftarrow$  maximum number of malicious processes,  $f < \lceil n/4 \rceil$ ;

(1) Each process executes the following  $f + 1$  phases, where  $f < n/4$ :

(1a) **for**  $phase = 1$  **to**  $f + 1$  **do**

(1b)   Execute the following Round 1 actions:                      // actions in round one of each phase

(1c)       **broadcast**  $v$  to all processes;

(1d)       **await** value  $v_j$  from each process  $P_j$ ;

(1e)        $majority \leftarrow$  the value among the  $v_j$  that occurs  $> n/2$  times (default if no maj.);

(1f)        $mult \leftarrow$  number of times that  $majority$  occurs;

(1g)   Execute the following Round 2 actions:                      // actions in round two of each phase

(1h)       **if**  $i = phase$  **then** // only the phase leader executes this send step

(1i)        **broadcast**  $majority$  to all processes;

(1j)        **receive**  $tiebreaker$  from  $P_{phase}$  (default value if nothing is received);

(1k)        **if**  $mult > n/2 + f$  **then**

(1l)            $v \leftarrow majority$ ;

(1m)        **else**  $v \leftarrow tiebreaker$ ;

(1n)        **if**  $phase = f + 1$  **then**

(1o)           output decision value  $v$ .

# The Phase King Algorithm

- $(f + 1)$  rounds,  $(f + 1)[(n - 1)(n + 1)]$  messages, and can tolerate up to  $f < \lceil n/4 \rceil$  malicious processes

## Correctness Argument

- Among  $f + 1$  rounds, at least one round  $k$  where phase-king is non-malicious.
- In round  $k$ , all non-malicious processes  $P_i$  and  $P_j$  will have same estimate of consensus value as  $P_k$  does.
  - ▶  $P_i$  and  $P_j$  use their own majority values (Hint:  $\implies P_i$ 's *mult*  $> n/2 + f$ )
  - ▶  $P_i$  uses its majority value;  $P_j$  uses phase-king's tie-breaker value.
  - ▶  $P_i$  and  $P_j$  use the phase-king's tie-breaker value. (Hint: In the round in which  $P_k$  is non-malicious, it sends same value to  $P_i$  and  $P_j$ )

In all 3 cases, argue that  $P_i$  and  $P_j$  end up with same value as estimate

- If all non-malicious processes have the value  $x$  at the start of a round, they will continue to have  $x$  as the consensus value at the end of the round.



## The Berman-Garray algorithm: Code for $p_i$ , $0 \leq i \leq n-1$

$\text{pref}[i] = x_i$ ,  $\text{pref}[j = v_\perp]$  for any  $0 \leq j \leq n-1$ ,  $j \neq i$

Round  $2k-1$ ,  $1 \leq k \leq f+1$ :

broadcast( $\text{pref}[i]$ )

receive  $v_j$  from  $p_j$

$\forall j: 0 \leq j \leq n-1, j \neq i$   $\text{pref}[j] := v_j$

$\text{maj} := \text{majority}\{\text{pref}[0], \dots, \text{pref}[n-1]\}$

$\text{mult} = \text{multiplicity of } \text{maj}$  /\*#procs that voted for  $\text{maj}$ \*/

Round  $2k$ ,  $1 \leq k \leq f+1$ :

if  $i = k$  then broadcast( $\text{maj}$ ) //  $p_i$  is the king

receive( $\text{king\_maj}$ ) from  $p_k$

if ( $\text{mult} > \frac{n}{2} + f$ )

then  $\text{pref}[i] := \text{maj}$

else  $\text{pref}[i] := \text{king\_maj}$

if ( $k = f+1$ ) then  $\text{decision} := \text{pref}[i]$

### Lemma 3

If all nonfaulty processes prefer  $v$  at the beginning of phase  $k$ , then they all prefer  $v$  at the end of phase  $k$ , for all  $1 \leq k \leq f + 1$ .

Proof sketch:

- Each  $p_i$  receives at least  $n - f$  copies of  $v$  (including its own) in the first round of phase  $k$ .
- Because of the assumption  
$$n > 4f \Rightarrow n/2 > 2f \Rightarrow n > n/2 + 2f \Rightarrow n - f > n/2 + f.$$
- Thus, all nonfaulty  $p_i$ s will prefer  $v$  at the end of phase  $k$ .

Validity follows by lemma 3: If all nonfaulty  $p_i$ s start with  $v$ , they continue to prefer  $v$  throughout the phase.

- Observation: There are at most  $f$  faulty  $p_i$ s, and  $f + 1$  phases, so at least one phase has a nonfaulty king.

### Lemma 4

Let  $g$  be a phase whose king  $p_g$  is nonfaulty. Then all nonfaulty  $p_i$ s finish phase  $g$  with the same preference.

Proof sketch:

- Case 1: Suppose all nonfaulty  $p_i$ s use king's majority for their preference. Since the king is nonfaulty it sends everyone the same value.

Proof sketch of lemma 4 continued:

- Case 2: Suppose that some nonfaulty  $p_j$  uses its own majority value  $v$  for its preference. This means that  $p_j$  has received more than  $n/2 + f$  votes for  $v$  in the first round of  $g$ . Thus, every processor, including the king  $p_g$ , has received more than  $n/2$  votes for  $v$  in the first round of  $g$ , and sets its majority value to  $v$ .
- Agreement follows by lemma 4: At phase  $g + 1$  all nonfaulty  $p_i$ s start with the same preference and by lemma 3 this agreement persists.